

Testing scientific code, Part II

Because you're worth it

Lisa Schwetlick and Pietro Berkes

Testing patterns

What a good test looks like

- ▶ What does a good test look like? What should I test?
- ▶ **Good:**
 - ▶ Short and quick to execute
 - ▶ Easy to read
 - ▶ Exercise *one* thing
- ▶ **Bad:**
 - ▶ Relies on data files
 - ▶ Messes with “real-life” files, servers, databases

Basic structure of test

- ▶ A good test is divided in three parts:
 - ▶ **Given:** Put your system in the right state for testing
 - ▶ Create data, initialize parameters, define constants...
 - ▶ **When:** Execute the feature that you are testing
 - ▶ Typically one or two lines of code
 - ▶ **Then:** Compare outcomes with the expected ones
 - ▶ Define the expected result of the test
 - ▶ Set of *assertions* that check that the new state of your system matches your expectations

Test simple but general cases

- ▶ Start with simple, general case
 - ▶ Take a realistic scenario for your code, try to reduce it to a simple example
- ▶ Tests for 'lower' method of strings

```
def test_lower():  
    # Given  
    string = 'HeLlO wOrld'  
    expected = 'hello world'  
  
    # When  
    output = string.lower()  
  
    # Then  
    assert output == expected
```

Test special cases and boundary conditions

- ▶ Code often breaks in corner cases: empty lists, None, NaN, 0.0, lists with repeated elements, non-existing file, ...
- ▶ This often involves making design decision: respond to corner case with special behavior, or raise meaningful exception?

```
def test_lower_empty_string():  
    # Given  
    string = ''  
    expected = ''  
  
    # When  
    output = string.lower()  
  
    # Then  
    assert output == expected
```

- ▶ Other good corner cases for `string.lower()`:
 - ▶ 'do-nothing case': `string = 'hi'`
 - ▶ symbols: `string = '123 (!'`

Common testing pattern

- ▶ Often these cases are collected in a single test:

```
def test_lower():  
    # Given  
    # Each test case is a tuple of (input, expected_result)  
    test_cases = [('HeLlO wOrld', 'hello world'),  
                  ('hi', 'hi'),  
                  ('123 ([?', '123 ([?'),  
                  ('', '')]  
  
    for string, expected in test_cases:  
        # When  
        output = string.lower()  
        # Then  
        assert output == expected
```

Parametrize

- ▶ Sometimes you want to run the same test multiple times with different values
- ▶ Option 1: for loop in your test
- ▶ Option 2: parametrize

```
@pytest.mark.parametrize("a", [1, 2, 3, 4])  
def test_addition_increases(a):  
    assert 5+a>a
```


Parametrize

- ▶ ... is also useful when you want to test different cases and their outcomes!



```
import pytest

@pytest.mark.parametrize("string, expected",
                        [('HeLlO wOrld', 'hello world'),
                        ('hi', 'hi'),
                        ('', '')])

def test_lower(string, expected):
    # When
    output = string.lower()
    # Then
    assert output == expected
```

Excursion: Logistic Map




- Sometimes used as a simple model for population growth



Excursion: Logistic Map

- ▶ x_0 should be between 0 and 1
- ▶ $f(x) = r * x * (1 - x)$
- ▶ Iterated function: $f(x_0) = x_1 \rightarrow f(x_1) = x_2 \rightarrow f(x_2) = x_3$



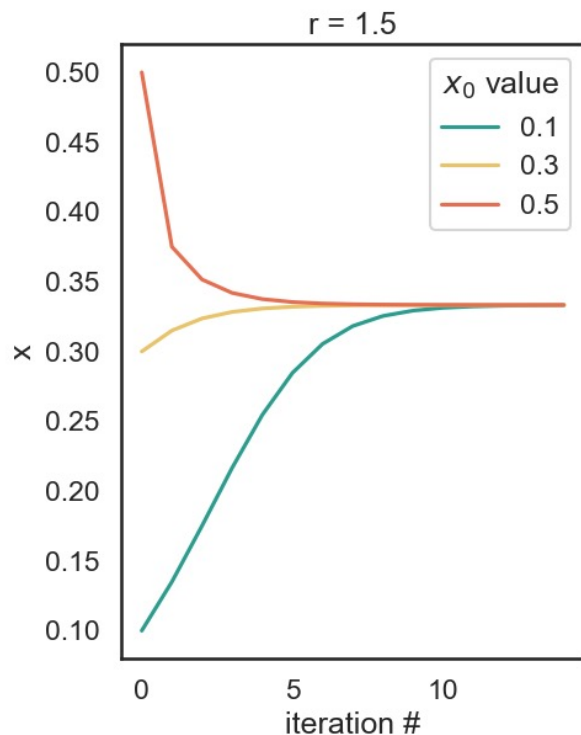
Logistic Function



Excursion: Logistic Map

- ▶ Looking at these plots, what could you test?

Logistic Function



Hands-on!

First fork the repo <https://github.com/aspp-latam/2023-mexico-testing-project> on GitHub and clone your own copy!

READ THE INSTRUCTIONS IN THE README.md!

Here's a summary:

a) Implement the logistic map $f(x)=r*x*(1-x)$. Use `@pytest.mark.parametrize` to test the function for the following cases:

- ▶ `x=0.1, r=2.2 => f(x, r)=0.198`
- ▶ `x=0.2, r=3.4 => f(x, r)=0.544`
- ▶ `x=0.75, r=1.7 => f(x, r)=0.31875`

b) Implement the function `iterate_f` that runs `f` for `it` iterations, each time passing the result back into `f`. Use `@pytest.mark.parametrize` to test the function for the following cases:

- ▶ `x=0.1, r=2.2, it=1`
`=> iterate_f(it, x, r)=[0.198]`
- ▶ `x=0.2, r=3.4, it=4`
`=> iterate_f(it, x, r)=[0.544, 0.843418, 0.449019, 0.841163]`
- ▶ `x=0.75, r=1.7, it=2`
`=> iterate_f(it, x, r)=[0.31875, 0.369152]`

c) Import and use the `plot_trajectory` function from the `plot_logfun` module to look at the trajectories generated by your code. Try with values `r<3`, `r>4`, and `3<r<4` to get an intuition for how the function behaves differently with different parameters.

Marking tests (xfail)

- ▶ Aside from `parametrize`, there are some other built in markers
- ▶ Sometimes you have a test that fails, but for good reason or you just want to deal with it later...
- ▶ Expected failure (`xfail`)
- ▶ Outputs an “x” (or “X”) in place of the “.”

```
@pytest.mark.xfail
def test_something():
    ...
```

Marking tests (skip)

- ▶ It is also possible to skip tests
- ▶ Useful when the feature doesn't exist yet or the test is very slow

```
@pytest.mark.skip(reason="functionality not yet  
implemented")  
def test_something():  
    ...
```

Marking tests with custom markers

- ▶ If you have lots of tests, you can categorize them with your own markers
 - ▶ although for custom mark names you need to register the marks “pytest.ini”
 - ▶ <https://docs.pytest.org/en/7.1.x/example/markers.html#registering-markers>
- ▶ Example:
 - ▶ Smoke tests check for really basic failure: run these frequently
 - ▶ Other tests may be many or too slow to run every time and test for more edge cases

```
@pytest.mark.smoke  
def test_something_basic():  
    ...
```


```
> pytest -m smoke  
> pytest -m "smoke and not slow"
```


Strategies for testing scientific code

Strategies for testing learning algorithms


- ▶ Learning algorithms can get stuck in local maxima, the solution for general cases might not be known (e.g., unsupervised learning)
- ▶ Turn your validation cases into tests
- ▶ Stability tests:
 - ▶ Start from final solution; verify that the algorithm stays there
 - ▶ Start from solution and add a small amount of noise to the parameters; verify that the algorithm converges back to the solution
- ▶ Parameter Recovery: Generate synthetic data from the model with known parameters, then test that the code can learn the parameters back

Learning algorithms fit the parameters of a model to observed data



Generate synthetic data from the model to test the learning algorithm by recovering the parameters

I) Fix initial parameters





$$a^* = 0.5$$

$$b^* = -1.3$$


Generate synthetic data from the model to test the learning algorithm by recovering the parameters

1) Fix initial parameters





 $y = a^* x + b^*$
+ noise

2) Generate synthetic data




Generate synthetic data from the model to test the learning algorithm by recovering the parameters

1) Fix initial parameters




2) Generate synthetic data


$$y = a^* x + b^* + \text{noise}$$



3) Run the algorithm




$$y = ax + b + \text{noise}$$




Generate synthetic data from the model to test the learning algorithm by recovering the parameters

1) Fix initial parameters




2) Generate synthetic data

$$y = a^* x + b^* + \text{noise}$$




4) Compare

$$a = 0.5098$$
$$b = -1.287$$



3) Run the algorithm

$$y = ax + b + \text{noise}$$



Other common cases

- ▶ Test general routines with specific ones
 - ▶ Example: `test polynomial_expansion(data, degree)`
with `quadratic_expansion(data)`
- ▶ Test optimized routines with brute-force approaches
 - ▶ Example: test function computing analytical derivative with numerical derivative

Randomness in Testing

- ▶ Using randomness in testing can be useful
 - ▶ For confirming generalizability and stability
 - ▶ For finding corner cases or numerical problems
 - ▶ Using Random/Sampled input data to test whether the result is as expected

```
def test_something():  
    for _ in range(10):  
        r = np.random.rand()  
        assert my_random_function(r)
```



Random Seeds and Reproducibility

- ▶ When running tests that involve randomness and some test doesn't pass it is vital to be able to reproduce that test exactly!
- ▶ Computers produce pseudo-random numbers: setting a seed resets the basis for the random number generator
- ▶ This is essential for reproducibility
- ▶ At a minimum, you should manually set the seed for your random test

```
SEED = 42
```

```
random_state = np.random.RandomState(SEED)  
random_state.rand()
```

Hands On!

READ THE INSTRUCTIONS IN THE README.md!

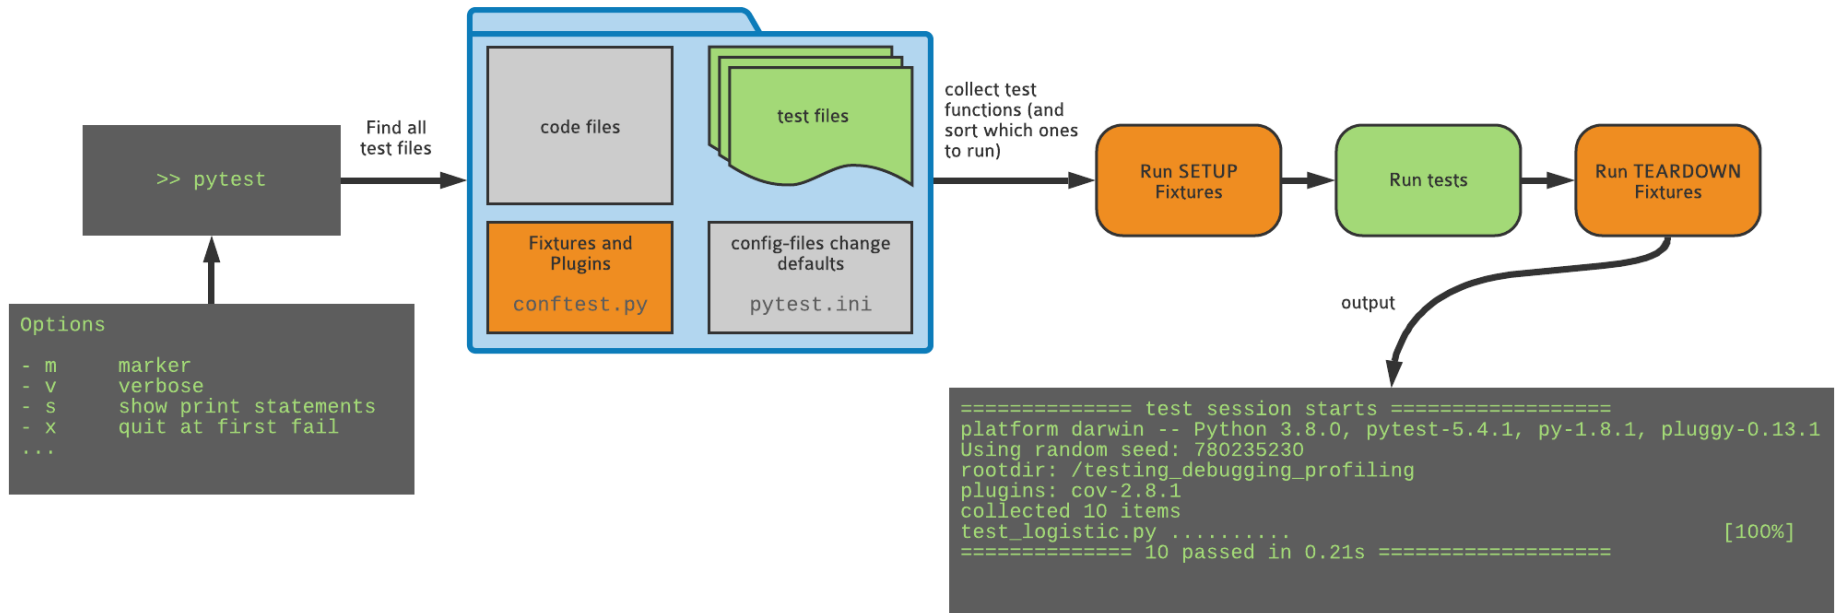
a) Write a randomized test that checks that, for $r=1.5$, any random starting points converge to the attractor $f(x, r) = 1/3$.



A Pytest Solution

- ▶ This is not so prominent in the docs, because non-scientific coding uses random testing more rarely
- ▶ In scientific coding, when you deal with randomness it is very relevant
- ▶ What do we want?
 - ▶ For each (random) test there should be a seed
 - ▶ For each run of the test, the seed should be different
 - ▶ That seed should be printed with the test result
 - ▶ It needs to be possible to explicitly run the test again with that seed!

Pytest



Fixtures (minimal solution)

- ▶ Fixtures are functions that are run before the tests are executed

```
import numpy as np
import pytest
```

```
# set the random seed for once here
SEED = np.random.randint(0, 2**31)
```

```
@pytest.fixture
def random_state():
    print(f'Using seed {SEED}')
    random_state = np.random.RandomState(SEED)
    return random_state
```

```
def test_something(random_state):
    random_state.rand()
```

Hands On!


READ THE INSTRUCTIONS IN THE README.md!

- a) Write a randomized test that checks that, for $r=1.5$, any random starting points converge to the attractor $f(x, r) = 1/3$.
- b) Add a fixture at the top of your test file, that lets you print the seed to the console.

```
[L]$ pytest
===== test session starts =====
platform darwin -- Python 3.8.0, pytest-5.4.1, py-1.8.1, pluggy-0.13.1
Using random seed: 892358865
```

Fixtures (real solution)

- ▶ `conftest.py` is a magical file! (don't import it!)
- ▶ `conftest.py` can be used to define custom behavior or plugins. Fixtures can also be defined here, so that they can be used by all tests.
- ▶ See the file `conftest_example.py` in the repo you forked. If you rename it the functions defined there select a seed for each test and allow you to pass a seed on the commandline using `--seed 123`




Hands On!

READ THE INSTRUCTIONS IN THE README.md!

- a) Write a randomized test that checks that, for $r=1.5$, any random starting points converge to the attractor $f(x, r) = 1/3$.
- b) Add a fixture at the top of your test file, that lets you print the seed to the console.
- c) Add a `conftest.py` file to set a random seed before each run and make possible failures reproducible
- d) Check that the console output of pytest now includes the seed!


```
[L]$ pytest  
  
===== test session starts =====  
platform darwin -- Python 3.8.0, pytest-5.4.1, py-1.8.1, pluggy-0.13.1  
Using random seed: 892358865
```

Excursion: Logistic Equation




- ▶ Between $r=3$ and $r=4$ the logistic map has a range of behaviors
- ▶ Periodic vs. chaotic

Excursion: Logistic Equation



Excursion: Logistic Equation

Logistic Function



- ▶ Sensitive Dependence on Initial Conditions (SDIC)
- ▶ Even seeds that are very close, quickly find completely different itineraries
- ▶ Butterfly effect



Hands on!

Some r values for $3 < r < 4$ have some interesting properties: a chaotic trajectory neither diverges nor converges.

a) Use the `plot_bifurcation` function from the `plot_logfun` module using your implementation of `f` and `iterate_f` to look at the bifurcation diagram. The function generates an output image, `bifurcation_diagram.png`

b) Write a test that checks for chaotic behavior when $r=3.8$. Run the logistic map for 100000 iterations and verify the conditions for chaotic behavior:

- 1) The function is deterministic: *this does not need to be tested in this case*
- 2) Orbits must be bounded: check that all values are between 0 and 1
- 3) Orbits must be aperiodic: check that the last 1000 values are all different
- 4) Sensitive dependence on initial conditions: *this is the bonus exercise (in readme)*

The test should check conditions 2) and 3)!

Testing is good for your self-esteem

- ▶ Immediately: Always be confident that your results are correct, whether your approach works or not
- ▶ In the future: **save your future-self some trouble!**
- ▶ If you are left thinking “it’s cool but I cannot test *my* code because XYZ”, talk to us during the week and we’ll show you how to do it ;-)

You, in 2023 \Rightarrow You, in 2024



Thank you!

Example: eigenvector decomposition

- ▶ Consider the function `values, vectors = eigen(matrix)`
- ▶ Test with simple but general cases:
 - ▶ use full matrices for which you know the exact solution (from a table or computed by hand)
- ▶ Test general routine with specific ones:
 - ▶ use the analytical solution for 2x2 matrices
- ▶ Generate data from the model:
 - ▶ generate random eigenvalues, random eigenvector; construct the matrix; then check that the function returns the correct eigenvalues and -vectors
- ▶ Test with boundary cases:
 - ▶ test with diagonal matrix: is the algorithm stable?
 - ▶ test with a singular matrix: is the algorithm robust? Does it raise appropriate error when it fails?